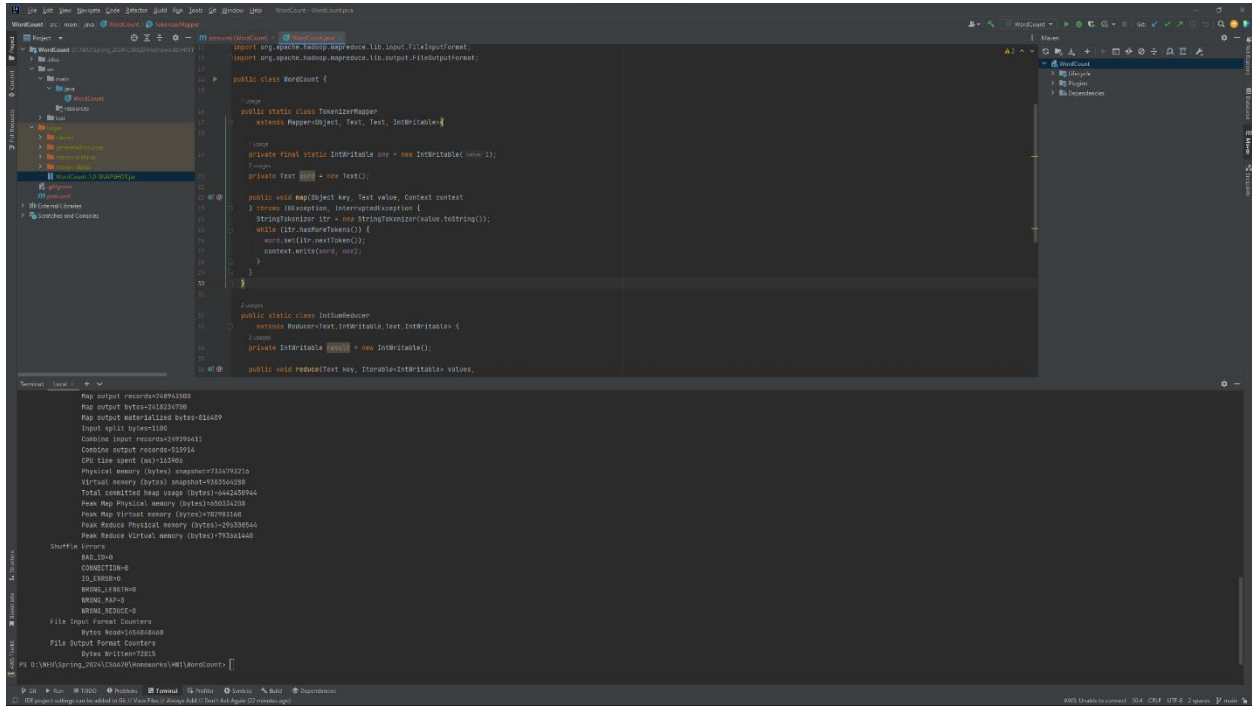


Name: Khoa Cu Dang (Cody) Cao

Course: CS6620 Spring 2024

Homework 1 Report

Local IDE Results



AWS Results

EMR Cluster Configuration

Instance groups (2) Info					
With the instance groups configuration, each node type consists of the same instance type and the same purchasing option for instances: On-Demand or Spot.					
Find instances by status		Find resources by ID or type; or search for text within loaded results			
Type and name	ID	Status	Instances	Purchasing option and p...	
Primary		Running	1	On-Demand	
m1.medium	ci-09784862TMF6PTC...	Running	-	On-Demand (\$0.087/hr)	
Core		Running	2	On-Demand	
m1.medium	ci-01921213DG7GU0F...	Running	-	On-Demand (\$0.087/hr)	
m1.medium	ci-0095521KRCXNWY...	Running	-	On-Demand (\$0.087/hr)	

Step Execution Result

<input type="checkbox"/>	Step ID	Status	Name	Log files	Creation time (UTC-08:00)	Start time (UTC-08:00)	Elapsed time
<input type="checkbox"/>	s-0653764279X26LYLZ23Z	Completed	WordCount JAR	controller syslog stderr stdout	February 5, 2024 at 14:09	February 5, 2024 at 14:09	5 minutes, 46 seconds
Jar location s3://cs6240-hw1-cody/WordCount-1.0-SNAPSHOT.jar				Permissions -	Main class -		
Action on failure Continue				Argument WordCount s3://cs6240-hw1-cody/input/hw1.txt s3://cs6240-hw1-cody/output/			

Input file on S3

Amazon S3 > Buckets > cs6240-hw1-cody > input/

input/ Copy S3 URI

Objects Properties

Objects (1) info

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	hw1.txt	txt	January 29, 2024, 14:05:22 (UTC-08:00)	1.4 GB	Standard

Output Directory on S3

Amazon S3 > Buckets > cs6240-hw1-cody > output/

output/ Copy S3 URI

Objects Properties

Objects (4) info

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	_SUCCESS	-	February 5, 2024, 14:15:30 (UTC-08:00)	0 B	Standard
<input type="checkbox"/>	part-00000	-	February 5, 2024, 14:15:25 (UTC-08:00)	23.3 KB	Standard
<input type="checkbox"/>	part-00001	-	February 5, 2024, 14:15:25 (UTC-08:00)	23.8 KB	Standard
<input type="checkbox"/>	part-00002	-	February 5, 2024, 14:15:28 (UTC-08:00)	24.0 KB	Standard

JAR file location on S3

Amazon S3 > Buckets > cs6240-hw1-cody

cs6240-hw1-cody info

Objects Properties Permissions Metrics Management Access Points

Objects (5) info

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	bootstrap_debug.sh	sh	February 4, 2024, 23:35:10 (UTC-08:00)	22.0 B	Standard
<input type="checkbox"/>	input/	Folder	-	-	-
<input type="checkbox"/>	logs/	Folder	-	-	-
<input type="checkbox"/>	output/	Folder	-	-	-
<input type="checkbox"/>	WordCount-1.0-SNAPSHOT.jar	jar	February 5, 2024, 13:58:56 (UTC-08:00)	4.6 KB	Standard

Last lines of syslog.txt

```
syslog.txt
File Edit View
2024-02-05 22:15:21,950 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 20%
2024-02-05 22:15:22,953 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 21%
2024-02-05 22:15:23,956 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 21%
2024-02-05 22:15:25,961 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 67%
2024-02-05 22:15:28,973 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2024-02-05 22:15:30,985 INFO org.apache.hadoop.mapreduce.Job (main): Job job_1707169248635_0002 completed successfully
2024-02-05 22:15:31,228 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 61

File System Counters
  FILE: Number of bytes read=6025225
  FILE: Number of bytes written=12766216
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=2178
  HDFS: Number of bytes written=0
  HDFS: Number of read operations=22
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=0
  HDFS: Number of bytes read erasure-coded=0
  S3: Number of bytes read=1454182711
  S3: Number of bytes written=72815
  S3: Number of read operations=0
  S3: Number of large read operations=0
  S3: Number of write operations=0

Job Counters
  Killed map tasks=1
  Killed reduce tasks=1
  Launched map tasks=22
  Launched reduce tasks=3
  Data-local map tasks=22
  Total time spent by all maps in occupied slots (ms)=20207016
  Total time spent by all reduces in occupied slots (ms)=6692032
  Total time spent by all map tasks (ms)=841959
  Total time spent by all reduce tasks (ms)=209126
  Total vcore-milliseconds taken by all map tasks=841959
  Total vcore-milliseconds taken by all reduce tasks=209126
  Total megabyte-milliseconds taken by all map tasks=646624512
  Total megabyte-milliseconds taken by all reduce tasks=214145024

Map-Reduce Framework
  Map input records=21907700
  Map output records=248943500
  Map output bytes=2418234700
  Map output materialized bytes=1916143
  Input split bytes=2178
  Combine input records=248943500
  Combine output records=232012
  Reduce input groups=5273
  Reduce shuffle bytes=1916143
  Reduce input records=232012
  Reduce output records=5273
  Spilled Records=696036
  Shuffled Maps =66
  Failed Shuffles=0
  Merged Map outputs=66
  GC time elapsed (ms)=14964
  CPU time spent (ms)=410400
  Physical memory (bytes) snapshot=11395403776
  Virtual memory (bytes) snapshot=58716188672
  Total committed heap usage (bytes)=6135898112
  Peak Map Physical memory (bytes)=490098688
  Peak Map Virtual memory (bytes)=2315685888
  Peak Reduce Physical memory (bytes)=279863296
  Peak Reduce Virtual memory (bytes)=2599497728

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=1454182711

File Output Format Counters
  Bytes Written=72815
```