

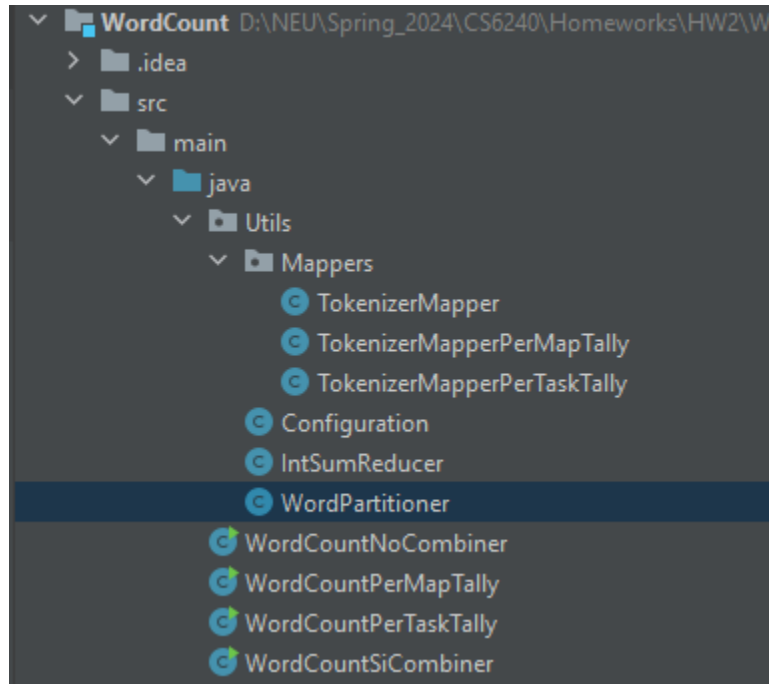
Name: Khoa Cu Dang (Cody) Cao

Course: CS6620 Spring 2024

Homework 2 Report

I. Souce Code

For code reuse purposes, I separated components into their own classes (Mapper, Reducer, Partitioner) then have different main classes utilize them as needed.



Common Utility Classes

Configuration: This class stores common configuration variables (list of acceptable starting characters, number of reducers being used)

```
8 usages  khoacao-cddk
7      public class Configuration {
      6 usages
8          public static final Set<Character> VALID_START_WORDS =
9              new HashSet<>(Arrays.asList('m', 'n', 'o', 'p', 'q', 'M', 'N', 'O', 'P', 'Q'));
      14 usages
10         public static final int NUM_REDUCER = 5;
11     }
12
```

IntSumReducer: This class is used for Reducer and in-mapper Combiner tasks.

```
/**
 * Reducer class - used to reduce/combine results
 */
9 usages  khoacao-ccdk
public class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    2 usages
    private IntWritable result = new IntWritable();

    khoacao-ccdk
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
    ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

WordPartitioner: This class is used to distribute key-value pairs emitted by mappers to reducers based on the word's starting character. Each reducer will handle a set of words that start with a specific character (both lowercase and uppercase).

```
9  /**
10  * WordPartitioner class - used to partition records to multiple reducers
11  */
12  8 usages khoacao-ccdk
13  public class WordPartitioner extends Partitioner<Text, IntWritable> {
14
15  /**
16  * Partitioning words into 5 different reducers. Each of them handles each specific letter
17  * @param word
18  * @param count
19  * @param i
20  * @return
21  */
22  no usages khoacao-ccdk
23  @Override
24  public int getPartition(Text word, IntWritable count, int i) {
25      char firstChar = word.toString().charAt(0);
26      switch (firstChar) {
27          case('m'):
28          case('M'):
29              return 0 % NUM_REDUCER;
30
31          case('n'):
32          case('N'):
33              return 1 % NUM_REDUCER;
34
35          case('o'):
36          case('O'):
37              return 2 % NUM_REDUCER;
38
39          case('p'):
40          case('P'):
41              return 3 & NUM_REDUCER;
42
43          default:
44              return 4 & NUM_REDUCER;
45      }
46  }
```

Mapper Classes

There are three main parameters for the `map()` function within Mapper classes:

- Key – This parameter is often not used and/or set to null since in mapper phase, the focus is on processing the values. This variable is also not mentioned as being used by Hadoop documentation.
- Value – From the document of Hadoop, this Text value contains a line of the file. This then gets broken down further into separate words (separated by space) using `StringTokenizer` class. Thus, each call of the `map()` function will handle one line of the input file.
- Context: This object allows mappers to interact with the rest of the Hadoop system, as explained by a [StackOverflow post](#). This explains the `context.write()` function being used to emit records.

TokenizerMapper: This is the default implementation of the WordCount example given by Hadoop. I went ahead and added an additional condition that the words being processed must start with a valid alphabetic character in order to be emitted.

```
11  /**
12   * Mapper class - Filtering "real" words and emit them
13   */
14  public class TokenizerMapper
15      extends Mapper<Object, Text, Text, IntWritable> {
16      private final static IntWritable one = new IntWritable(1);
17      private Text word = new Text();
18
19      public void map(Object key, Text value, Context context
20      ) throws IOException, InterruptedException {
21          StringTokenizer itr = new StringTokenizer(value.toString());
22          while (itr.hasMoreTokens()) {
23              word.set(itr.nextToken());
24              char firstChar = word.toString().charAt(0);
25              //Only emits the word if it starts with a valid character
26              if(VALID_START_WORDS.contains(firstChar)) {
27                  context.write(word, one);
28              }
29          }
30      }
31  }
32
```

TokenizerMapperPerMapTally: With this per map tally implementation, the emitting record will be held off until the map has gone through every word in the input line. Aggregation is also performed using `HashMap` in case a word appears more than once in a single line.

```
14  /**
15   * TokenizerMapperPerMapTally class - Filtering "real" words and emit them
16   *
17   * A HashMap is used to aggregate the count of every word iterated throughout the map function lifecycle
18   */
19  public class TokenizerMapperPerMapTally
20      extends Mapper<Object, Text, Text, IntWritable> {
21      private final static IntWritable one = new IntWritable(1);
22      private Text word = new Text();
23
24      @Override
25      public void map(Object key, Text value, Context context
26          ) throws IOException, InterruptedException {
27          Map<Text, IntWritable> countMap = new HashMap<>();
28
29          StringTokenizer itr = new StringTokenizer(value.toString());
30          while (itr.hasMoreTokens()) {
31              word.set(itr.nextToken());
32              char firstChar = word.toString().charAt(0);
33              //Only emits the word if it starts with a valid character
34              if(VALID_START_WORDS.contains(firstChar)) {
35                  IntWritable currCount = countMap.getOrDefault(word, new IntWritable(0));
36
37                  //Increment count
38                  currCount.set(currCount.get() + 1);
39                  countMap.put(new Text(word.toString()), currCount);
40              }
41          }
42
43          //Emits the final counter values after every record has been iterated
44          for(Entry<Text, IntWritable> record : countMap.entrySet()) {
45              context.write(record.getKey(), record.getValue());
46          }
47      }
48  }
```

TokenizerMapperPerTaskTally: With this mapper implementation, I am having the aggregation at the mapper/input split level. Records are being held from emitting until the mapper has finished processing the whole input split. The records are aggregated using HashMap, then emitted using the cleanup() method of the mapper.

```
14  /**
15   * TokenizerMapperPerTaskTally class - Filtering "real" words and emit them
16   *
17   * A HashMap is used to aggregate the count of every word iterated throughout the instance's lifecycle
18   */
19  public class TokenizerMapperPerTaskTally
20      extends Mapper<Object, Text, Text, IntWritable> {
21      private final static IntWritable one = new IntWritable(1);
22      private Text word = new Text();
23      private Map<Text, IntWritable> countMap = new HashMap<>();
24
25      @Override
26      public void map(Object key, Text value, Context context
27          ) throws IOException, InterruptedException {
28          StringTokenizer itr = new StringTokenizer(value.toString());
29          while (itr.hasMoreTokens()) {
30              word.set(itr.nextToken());
31              char firstChar = word.toString().charAt(0);
32              //Only emits the word if it starts with a valid character
33              if(VALID_START_WORDS.contains(firstChar)) {
34                  IntWritable currCount = countMap.getOrDefault(word, new IntWritable(0));
35
36                  //Increment count
37                  currCount.set(currCount.get() + 1);
38                  countMap.put(new Text(word.toString()), currCount);
39              }
40          }
41
42      @Override
43      protected void cleanup(Context context) throws IOException, InterruptedException {
44          // Emits the final counter values after all records have been processed in the input split
45          for (Entry<Text, IntWritable> record : countMap.entrySet()) {
46              context.write(record.getKey(), record.getValue());
47          }
48      }
49  }
```

Job Classes

WordCountNoCombiner

```
16 ▶ public class WordCountNoCombiner {
17   @ khoacao-ccdk *
18   public static void main(String[] args) throws Exception {
19     Configuration conf = new Configuration();
20     Job job = Job.getInstance(conf, jobName: "word count - no combiner");
21     job.setJarByClass(WordCountNoCombiner.class);
22     job.setMapperClass(TokenizerMapper.class);
23     job.setPartitionerClass(WordPartitioner.class);
24     //Setting number of reducer in accordance to the number of words
25     job.setReducerClass(IntSumReducer.class);
26     job.setNumReduceTasks(NUM_REDUCER);
27     job.setOutputKeyClass(Text.class);
28     job.setOutputValueClass(IntWritable.class);
29     FileInputFormat.addInputPath(job, new Path(args[0]));
30     FileOutputFormat.setOutputPath(job, new Path(args[1]));
31     System.exit(job.waitForCompletion(verbose: true) ? 0 : 1);
32   }
33 }
```

WordCountSiCombiner

```
16 ▶ public class WordCountSiCombiner {
17   @ khoacao-ccdk
18   public static void main(String[] args) throws Exception {
19     Configuration conf = new Configuration();
20     Job job = Job.getInstance(conf, jobName: "word count");
21     job.setJarByClass(WordCountSiCombiner.class);
22     job.setMapperClass(TokenizerMapper.class);
23     job.setCombinerClass(IntSumReducer.class);
24     job.setPartitionerClass(WordPartitioner.class);
25     //Setting number of reducer in accordance to the number of words
26     job.setReducerClass(IntSumReducer.class);
27     job.setNumReduceTasks(NUM_REDUCER);
28     job.setOutputKeyClass(Text.class);
29     job.setOutputValueClass(IntWritable.class);
30     FileInputFormat.addInputPath(job, new Path(args[0]));
31     FileOutputFormat.setOutputPath(job, new Path(args[1]));
32     System.exit(job.waitForCompletion(verbose: true) ? 0 : 1);
33   }
34 }
```

WordCountPerMapTally

```
14 ▶ public class WordCountPerMapTally {  
    khoacao-ccdk  
15 ▶ @ public static void main(String[] args) throws Exception {  
16     Configuration conf = new Configuration();  
17     Job job = Job.getInstance(conf, jobName: "word count");  
18     job.setJarByClass(WordCountPerMapTally.class);  
19     job.setMapperClass(TokenizerMapperPerMapTally.class);  
20     job.setPartitionerClass(WordPartitioner.class);  
21     //Setting number of reducer in accordance to the number of words  
22     job.setReducerClass(IntSumReducer.class);  
23     job.setNumReduceTasks(NUM_REDUCER);  
24     job.setOutputKeyClass(Text.class);  
25     job.setOutputValueClass(IntWritable.class);  
26     FileInputFormat.addInputPath(job, new Path(args[0]));  
27     FileOutputFormat.setOutputPath(job, new Path(args[1]));  
28     System.exit(job.waitForCompletion(verbose: true) ? 0 : 1);  
29 }  
30 }
```

WordCountPerTaskTally

```
15 ▶ public class WordCountPerTaskTally {  
    khoacao-ccdk  
16 ▶ @ public static void main(String[] args) throws Exception {  
17     Configuration conf = new Configuration();  
18     Job job = Job.getInstance(conf, jobName: "word count");  
19     job.setJarByClass(WordCountPerTaskTally.class);  
20     job.setMapperClass(TokenizerMapperPerTaskTally.class);  
21     job.setPartitionerClass(WordPartitioner.class);  
22     //Setting number of reducer in accordance to the number of words  
23     job.setReducerClass(IntSumReducer.class);  
24     job.setNumReduceTasks(NUM_REDUCER);  
25     job.setOutputKeyClass(Text.class);  
26     job.setOutputValueClass(IntWritable.class);  
27     FileInputFormat.addInputPath(job, new Path(args[0]));  
28     FileOutputFormat.setOutputPath(job, new Path(args[1]));  
29     System.exit(job.waitForCompletion(verbose: true) ? 0 : 1);  
30 }  
31 }
```


II. Performance Comparison

Record of 16 Steps being run on AWS. After the first 8 steps, I edited the cluster configuration to 10 core nodes and 1 master node.

Properties

Bootstrap actions

Instances (Hardware)

Steps

Applications

Configurations

Monitoring

Events

Tags (1)

Steps (16) [info](#)

Each step is a unit of work that contains instructions to manipulate data for processing by software installed on the cluster.

Concurrent steps: 1 [↗](#)

Filter steps by status

<input type="checkbox"/>	Step ID	Status	Name	Log files ↗	Creation time (UTC-08:00)	Start time (UTC-08:00)	Elapsed time
<input type="checkbox"/>	s-0937167FSCQUY7TMT	Completed	WordCount-PerTaskTally-Run2-11instances	controller syslog stderr stdout ↗	February 20, 2024 at 21:49	February 20, 2024 at 21:58	1 minute, 34 seconds
<input type="checkbox"/>	s-00845713324ZYQKR6M	Completed	WordCount-PerTaskTally-Run1-11instances	controller syslog stderr stdout ↗	February 20, 2024 at 21:49	February 20, 2024 at 21:57	1 minute, 24 seconds
<input type="checkbox"/>	s-01724783KHD42X11FFGK	Completed	WordCount-PerMapTally-Run2-11instances	controller syslog stderr stdout ↗	February 20, 2024 at 21:48	February 20, 2024 at 21:54	1 minute, 58 seconds
<input type="checkbox"/>	s-080491321M5PTWQ9ISNO	Completed	WordCount-PerMapTally-Run1-11instances	controller syslog stderr stdout ↗	February 20, 2024 at 21:48	February 20, 2024 at 21:52	1 minute, 56 seconds
<input type="checkbox"/>	s-0184572XVD48GNG8GCA	Completed	WordCount-SiCombiner-Run2-11instances	controller syslog stderr stdout ↗	February 20, 2024 at 21:47	February 20, 2024 at 21:51	1 minute, 40 seconds
<input type="checkbox"/>	s-05552152GYR8CDOT1B4M	Completed	WordCount-SiCombiner-Run1	controller syslog stderr stdout ↗	February 20, 2024 at 21:47	February 20, 2024 at 21:49	1 minute, 46 seconds
<input type="checkbox"/>	s-09225422IIS8S7R4KEZ	Completed	WordCount-NoCombiner-Run2-11instances	controller syslog stderr stdout ↗	February 20, 2024 at 21:46	February 20, 2024 at 21:47	1 minute, 38 seconds
<input type="checkbox"/>	s-09234033EZ577WW4ETH	Completed	WordCount-NoCombiner-Run1-11instances	controller syslog stderr stdout ↗	February 20, 2024 at 21:45	February 20, 2024 at 21:45	1 minute, 48 seconds
<input type="checkbox"/>	s-0029067DQ2HHICGH404P	Completed	WordCount-PerTaskTally-Run2	controller syslog stderr stdout ↗	February 20, 2024 at 21:25	February 20, 2024 at 21:25	1 minute, 50 seconds
<input type="checkbox"/>	s-053927524N19CNY0T4Z	Completed	WordCount-PerTaskTally-Run1	controller syslog stderr stdout ↗	February 20, 2024 at 21:20	February 20, 2024 at 21:21	1 minute, 52 seconds
<input type="checkbox"/>	s-01712145QUN7HHEXTF5	Completed	WordCount-PerMapTally-Run2	controller syslog stderr stdout ↗	February 20, 2024 at 21:14	February 20, 2024 at 21:15	2 minutes, 48 seconds
<input type="checkbox"/>	s-0209833SLQ3B49PU984X	Completed	WordCount-PerMapTally-Run1	controller syslog stderr stdout ↗	February 20, 2024 at 21:09	February 20, 2024 at 21:09	2 minutes, 42 seconds
<input type="checkbox"/>	s-013743531X4VT4KE7520	Completed	WordCount-SiCombiner-Run2	controller syslog stderr stdout ↗	February 20, 2024 at 21:05	February 20, 2024 at 21:05	2 minutes
<input type="checkbox"/>	s-0982400BUKRWTV0UHBE	Completed	WordCount-SiCombiner-Run1	controller syslog stderr stdout ↗	February 20, 2024 at 21:00	February 20, 2024 at 21:00	1 minute, 58 seconds
<input type="checkbox"/>	s-0553350WPX68EXEYIA9	Completed	WordCount-NoCombiner-Run2	controller syslog stderr stdout ↗	February 20, 2024 at 20:56	February 20, 2024 at 20:56	2 minutes, 38 seconds
<input type="checkbox"/>	s-081399523BW1H7OCBUXI	Completed	WordCount-NoCombiner-Run1	controller syslog stderr stdout ↗	February 20, 2024 at 20:14	February 20, 2024 at 20:14	2 minutes, 54 seconds

Events (50+)

info

Find events by source, severity or event type; or search for text within loaded results

Filter by a date and time range

resize

Clear filter

Showing 17 of 50+ events

Time (UTC-08:00)	Severity	Description	Event type	Event code	Source ID	Source type
February 20, 2024, 21:42	Info	The resizing operation for instance group ig-3UZTQKYP9ZUPX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) is complete. It now has an instance count of 10. The resize started at 2024-02-21 05:37 UTC and took 4 minutes to complete.	Instance Group State Change	-	ig-3UZTQKYP9ZUPX	Instance Group
February 20, 2024, 21:38	Info	The resizing operation for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) is complete. It now has an instance count of 0. The resize started at 2024-02-21 05:37 UTC and took 0 minutes to complete.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:38	Info	A resize for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) started at 2024-02-21 05:38 UTC. It is resizing from an instance count of 0 to 0.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:38	Info	A resize for instance group ig-3UZTQKYP9ZUPX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) started at 2024-02-21 05:38 UTC. It is resizing from an instance count of 5 to 10.	Instance Group State Change	-	ig-3UZTQKYP9ZUPX	Instance Group
February 20, 2024, 21:37	Info	A resize for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) was initiated by user at 2024-02-21 05:37 UTC.	Instance Group Status Notification	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:37	Info	A resize for instance group ig-3UZTQKYP9ZUPX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) was initiated by user at 2024-02-21 05:37 UTC.	Instance Group Status Notification	-	ig-3UZTQKYP9ZUPX	Instance Group
February 20, 2024, 21:37	Info	A resize for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) was initiated by user at 2024-02-21 05:37 UTC.	Instance Group Status Notification	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:37	Info	A resize for instance group ig-3UZTQKYP9ZUPX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) was initiated by user at 2024-02-21 05:37 UTC.	Instance Group Status Notification	-	ig-3UZTQKYP9ZUPX	Instance Group
February 20, 2024, 21:31	Info	The resizing operation for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) is complete. It now has an instance count of 0. The resize started at 2024-02-21 05:27 UTC and took 4 minutes to complete.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:27	Info	A resize for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) started at 2024-02-21 05:27 UTC. It is resizing from an instance count of 1 to 0.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:27	Info	The resizing operation for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) is complete. It now has an instance count of 1. The resize started at 2024-02-21 05:26 UTC and took 0 minutes to complete.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:26	Info	The resizing operation for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) is complete. It now has an instance count of 1. The resize started at 2024-02-21 05:26 UTC and took 0 minutes to complete.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:26	Info	A resize for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) started at 2024-02-21 05:26 UTC. It is resizing from an instance count of 1 to 1.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:23	Info	A resize for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) started at 2024-02-21 05:23 UTC. It is resizing from an instance count of 1 to 0.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:23	Info	The resizing operation for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) is complete. It now has an instance count of 1. The resize started at 2024-02-21 05:22 UTC and took 0 minutes to complete.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:22	Info	The resizing operation for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) is complete. It now has an instance count of 1. The resize started at 2024-02-21 05:22 UTC and took 0 minutes to complete.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group
February 20, 2024, 21:22	Info	A resize for instance group ig-16XPNPCHM9TXX in Amazon EMR cluster j-2YJ3B5G384X9Z (C56620-HW1-Cluster) started at 2024-02-21 05:22 UTC. It is resizing from an instance count of 1 to 1.	Instance Group State Change	-	ig-16XPNPCHM9TXX	Instance Group

Cluster Size	Main Class	Run	Start time	End time	Execution time (seconds)
6	NoCombiner	#1	2/21/2024 4:14:48	2/21/2024 4:17:32	164.00
6	NoCombiner	#2	2/21/2024 4:56:23	2/21/2024 4:58:54	151.00
6	SiCombiner	#1	2/21/2024 5:01:02	2/21/2024 5:02:52	110.00
6	SiCombiner	#2	2/21/2024 5:05:39	2/21/2024 5:07:32	113.00
6	PerMapTally	#1	2/21/2024 5:09:45	2/21/2024 5:12:21	156.00
6	PerMapTally	#2	2/21/2024 5:15:27	2/21/2024 5:18:08	161.00
6	PerTaskTally	#1	2/21/2024 5:21:06	2/21/2024 5:22:50	104.00
6	PerTaskTally	#2	2/21/2024 5:25:13	2/21/2024 5:26:55	102.00
11	NoCombiner	#1	2/21/2024 5:45:49	2/21/2024 5:47:29	100.00
11	NoCombiner	#2	2/21/2024 5:47:43	2/21/2024 5:49:14	91.00
11	SiCombiner	#1	2/21/2024 5:49:32	2/21/2024 5:51:06	94.00
11	SiCombiner	#2	2/21/2024 5:51:18	2/21/2024 5:52:51	93.00
11	PerMapTally	#1	2/21/2024 5:53:02	2/21/2024 5:54:52	110.00
11	PerMapTally	#2	2/21/2024 5:55:04	2/21/2024 5:56:55	111.00
11	PerTaskTally	#1	2/21/2024 5:57:06	2/21/2024 5:58:24	78.00
11	PerTaskTally	#2	2/21/2024 5:58:44	2/21/2024 6:00:04	80.00

- Do you believe the combiner was called at all in program SiCombiner?
- What difference did the use of a combiner make in SiCombiner compared to

NoCombiner?

Map-Reduce Framework	Map-Reduce Framework
Map input records=21907700	Map input records=21907700
Map output records=42842400	Map output records=42842400
Map output bytes=412253400	Map output bytes=412253400
Map output materialized bytes=23714267	Map output materialized bytes=186088
Input split bytes=2178	Input split bytes=2178
Combine input records=0	Combine input records=42842400
Combine output records=0	Combine output records=18678
Reduce input groups=849	Reduce input groups=849
Reduce shuffle bytes=23714267	Reduce shuffle bytes=186088
Reduce input records=42842400	Reduce input records=18678
Reduce output records=849	Reduce output records=849
Spilled Records=85684800	Spilled Records=37356
Shuffled Maps =110	Shuffled Maps =110
Failed Shuffles=0	Failed Shuffles=0
Merged Map outputs=110	Merged Map outputs=110
GC time elapsed (ms)=24799	GC time elapsed (ms)=23033
CPU time spent (ms)=246890	CPU time spent (ms)=222860
Physical memory (bytes) snapshot=12694753280	Physical memory (bytes) snapshot=11891810304
Virtual memory (bytes) snapshot=63892869120	Virtual memory (bytes) snapshot=63880028160
Total committed heap usage (bytes)=7256178688	Total committed heap usage (bytes)=6257664000
Peak Map Physical memory (bytes)=492720128	Peak Map Physical memory (bytes)=493527040
Peak Map Virtual memory (bytes)=2314612736	Peak Map Virtual memory (bytes)=2314940416
Peak Reduce Physical memory (bytes)=577753088	Peak Reduce Physical memory (bytes)=279035904
Peak Reduce Virtual memory (bytes)=2600955904	Peak Reduce Virtual memory (bytes)=2600951808
	Shuffle Errors

To answer both of the questions, we can look at the syslog for NoCombiner (left) and SiCombiner (right) run. When analyzing the syslog files, we can see that the combiner was definitely called in SiCombiner run due to:

- Combine input records (0 on NoCombiner vs 42842400 on SiCombiner)
- Combine output records (0 on NoCombiner vs 18678 on SiCombiner)
- Reduce input records (42842400 on NoCombiner vs 18678 on SiCombiner)

This proves that the combiner was used in the middle of the Map phase and the Reduce phase, which reduced the number of input records on the Reduce phase significantly by doing aggregation

before emitting records to the reduce phase. Additionally, having less input records means it took less time for the reduce tasks to run (321898ms on NoCombiner vs 175512ms on SiCombiner)

Total time spent by all map tasks (ms)=785307	Total time spent by all map tasks (ms)=711072
Total time spent by all reduce tasks (ms)=321898	Total time spent by all reduce tasks (ms)=175512
Total users milliseconds taken by all map tasks=785307	Total users milliseconds taken by all map tasks=711072

- **Was the local aggregation effective in PerMapTally compared to NoCombiner?**

Map output records=42842400	Map output records=40866300
Map output bytes=412253400	Map output bytes=396549900
Map output materialized bytes=23714267	Map output materialized bytes=27032788
Input split bytes=2178	Input split bytes=2178

Local aggregation in PerMapTally (right) seems to be a little bit better compared to NoCombiner (left)

- It has a slightly lower map output record (40866300 on PerMapTally vs 42842400 on NoCombiner)

This might be due to the fact that for each line, there is actually not many repetitive words. Thus, having a HashMap storing word count for only one line (PerMapTally) might not be an effective method when considering extra memory and additional runtime required to provision and operate an additional HashMap.

- **What differences do you see between PerMapTally and PerTaskTally? Try to explain the reasons.**

Map-Reduce Framework	Map-Reduce Framework
Map input records=21907700	Map input records=21907700
Map output records=40866300	Map output records=18678
Map output bytes=396549900	Map output bytes=229702
Map output materialized bytes=27032788	Map output materialized bytes=186088
Input split bytes=2178	Input split bytes=2178
Combine input records=0	Combine input records=0
Combine output records=0	Combine output records=0
Reduce input groups=849	Reduce input groups=849
Reduce shuffle bytes=27032788	Reduce shuffle bytes=186088
Reduce input records=40866300	Reduce input records=18678
Reduce output records=849	Reduce output records=849

When looking at the syslog of PerMapTally (left) vs PerTaskTally (right), there is a significant different in the Map out put record number (40866300 on PerMapTally vs 186088 on PerTaskTally), resulting in a significantly lower Reduce input records number. This can be because of the nature of the input file:

- Within a line of the input file, there is a fewer number of repetitive words. Thus, the HashMap provisioned in PerMapTally would have less value in aggregating the values.
- However, when we consider a split of the input file, there is a larger number of repetitive words. In this scenario, a HashMap provisioned at a mapper instance level (PerTaskTally) would help tremendously with aggregating records and reducing the number of output records.

- **Which one is better: SiCombiner or PerTaskTally? Briefly justify your answer.**

Map-Reduce Framework	Map-Reduce Framework
Map input records=21907700	Map input records=21907700
Map output records=18678	Map output records=42842400
Map output bytes=229702	Map output bytes=412253400
Map output materialized bytes=186088	Map output materialized bytes=186088
Input split bytes=2178	Input split bytes=2178
Combine input records=0	Combine input records=42842400
Combine output records=0	Combine output records=18678
Reduce input groups=849	Reduce input groups=849
Reduce shuffle bytes=186088	Reduce shuffle bytes=186088
Reduce input records=18678	Reduce input records=18678
Reduce output records=849	Reduce output records=849
Spilled Records=37356	Spilled Records=37356
Shuffled Maps =110	Shuffled Maps =110
Failed Shuffles=0	Failed Shuffles=0
Merged Map outputs=110	Merged Map outputs=110

Both results, PerTaskTally (left) and SiCombiner (right), while having slightly different aggregation strategies, result in similar number of input records in the reducer phase. Both approaches also maximize the use of parallel computing by either having an in-mapper aggregation using a HashMap or using a custom combiner that aggregating the output records of the mappers records locally. Thus, I would consider them as equally good.

- **NEW: Comparing the results for Configurations 1 and 2, do you believe this MapReduce program scales well to larger clusters? Briefly justify your answer**

When looking as the total Step runtime from EC2, we can see that the steps that ran with 11 instances tend to be around 20 seconds faster compared to those that ran with 6 instances. This does show that there is a benefit to running more cluster. However, when looking at other factors (CPU time spent, total time spent by all mappers) there is actually less difference between the two configuration. Which means that the program does not benefit from having more instances, thus does not scale particularly well with a larger cluster.